Sensitivity to Methodology

March 23, 2026

Contents

1 Overview 2

2 Data Quality and Missing Values 2

3 Variable Selection and Coefficient Diagnostics 4

4 Lag Structure and Cross-Correlation Analysis 7

5 Regime-Switching Fair Value (Realized Volatility vs VIX) 9

6 Rolling-Window Fair Value Model 12

7 Conclusions 14

1 Overview

This document examines the sensitivity of the main report's analytical results to alternative methodological choices. It begins with a data-quality assessment and then addresses the following analytical areas:

1. **Data quality and missing values** — a per-series inventory of missing observations, their causes, and the treatment applied throughout the analysis.
2. **Variable selection and coefficient diagnostics** — full OLS regression output for all candidate fair-value models, including coefficient estimates, standard errors, t-statistics, and information criteria (AIC/BIC) to support the model selection rationale.
3. **Lag structure and cross-correlation analysis** — autocorrelation and partial autocorrelation of CDX IG daily changes, together with cross-correlation functions against each candidate driver, to assess whether lagged regressors would improve model specification.
4. **Regime-switching static fair value (RV vs VIX)** — a sensitivity in which OLS coefficients depend on whether trailing realised volatility of the S&P index (from daily level returns) exceeds VIX by more than a fixed margin (percentage points on the spread), with a figure showing regime timing.
5. **Rolling-window fair value model** — whether allowing model coefficients to vary over time (rolling-window estimation) improves or degrades the fair-value specification relative to the static global-coefficient approach used in the main report.

2 Data Quality and Missing Values

Overview

The dataset comprises 1 304 daily observations across 20 time series, covering the period January 2018 to April 2023. Not all series are available for every trading day: holidays, differing market calendars (US vs. European), and data-provider coverage gaps introduce missing values.

Missing-Value Statistics

Table 1 summarises the number and proportion of missing observations for each series, sorted by the percentage of missing data.

Table 1: Missing observations per series.

Series	Missing	Available	% Missing
FR5BEI	1 112	192	85.3
ESTR	395	909	30.3
SOFR	57	1 247	4.4
UST10	55	1 249	4.2
UST2	55	1 249	4.2
US5BEI	55	1 249	4.2
CDX IG	54	1 250	4.1
CDX HY	54	1 250	4.1
S&P	46	1 258	3.5
ITRX IG	42	1 262	3.2
ITRX XO	42	1 262	3.2
VIX	39	1 265	3.0
DEM2	24	1 280	1.8
DEM10	23	1 281	1.8
ESTX600	19	1 285	1.5
CO	12	1 292	0.9
EURUSD	0	1 304	0.0
EURUSD3mBS	0	1 304	0.0
USYC	0	1 304	0.0
DEYC	0	1 304	0.0

Key Observations

- **FR5BEI (French 5yr inflation breakeven):** 85.3% of observations are missing. With only 192 usable data points out of 1 304, this series cannot support reliable statistical analysis. It is therefore **excluded from all analyses** in both the main report and this sensitivity document.
- **ESTR (Euro short-term rate):** 30.3% missing. ESTR was introduced by the ECB in October 2019, so observations before that date are unavailable. Any model that includes ESTR is effectively restricted to a roughly three-year subsample.
- **US rates, CDS indices, and equity series (3–4% missing):** these gaps align with US and European public holidays and occasional data-feed outages. Their impact on analysis is modest.
- **EURUSD, EURUSD3mBS, USYC, DEYC:** fully populated with no missing values.

Treatment

The following approach is applied consistently throughout the main report and this sensitivity document:

1. **Pairwise (listwise per analysis) deletion.** For each specific computation — whether a correlation, regression, or rolling-window estimate — only dates for which *all* required series are simultaneously available are used. This avoids imputing values for financial time series where interpolation could distort volatility clustering, jump dynamics, and cross-sectional relationships.
2. **No forward-filling or interpolation.** Given that the missing observations predominantly correspond to genuine market closures (holidays), carrying forward the previous day's close would artificially inject zero returns on non-trading days. This would understate realised volatility and bias correlation estimates toward zero.

3. **Impact on model selection.** Following the exclusion of FR5BEI, the InflationMacro candidate model is reduced to four regressors (ΔCO, $\Delta US5BEI$, $\Delta UST2$, $\Delta DEM2$). The high proportion of missing values in ESTR (30.3%) also materially reduces the effective sample for any model that includes it: the FXBasisPolicy specification retains only 834 observations after differencing, limiting the reliability of its rolling-window estimates.

3 Variable Selection and Coefficient Diagnostics

This section presents the full OLS regression diagnostics for all four candidate CDX IG fair-value models estimated on the training set (first 80% of the sample). For each model we report: coefficient estimates, standard errors, t-statistics, and model-level information criteria (AIC, BIC) that penalise over-parameterisation.

Model-level summary

Table 2 compares the four candidate specifications on key fit and selection metrics. Lower AIC and BIC values indicate a more parsimonious fit.

Table 2: Model-level diagnostics (training set).

Model	n	k	R^2	Adj. R^2	AIC	BIC
CrossMarket_d	916	6	0.516	0.514	4 015	4 044
RatesVolEquity_d	952	6	0.485	0.483	4 201	4 231
FXBasisPolicy_d	667	6	0.057	0.049	3 548	3 575
InflationMacro_d	940	5	0.146	0.142	4 632	4 657

Observations:

- On in-sample R^2 alone, CrossMarket_d ranks highest (0.516). However, this model includes iTraxx IG as a regressor—an index that is economically near-collinear with CDX IG—making it partially tautological. It is therefore not selected as the primary fair-value specification despite its in-sample fit.
- RatesVolEquity_d achieves the highest *out-of-sample* R^2 (0.698, as reported in the main document) despite a lower in-sample R^2, indicating better generalisation and less overfitting.
- AIC and BIC are not directly comparable across models with different sample sizes (FXBasisPolicy_d uses a shorter series due to ESTR data availability). Among the three models sharing comparable sample sizes ($\sim$ 920–950), CrossMarket_d has the lowest AIC/BIC, followed by RatesVolEquity_d, with InflationMacro_d substantially higher.
- FXBasisPolicy_d has very low explanatory power ($R^2 = 0.057$), confirming that FX and short-rate variables alone are insufficient to explain IG credit spread movements.

Coefficient detail: RatesVolEquity_d (selected model)

Table 3 presents the coefficient estimates for the selected specification.

Table 3: OLS coefficients — RatesVolEquity_d (training set, $n = 952$).

Variable	Coefficient	Std. Error	t-statistic	Sig. (5%)
Intercept	+0.077	0.071	+1.08	No
ΔVIX	+0.268	0.053	+5.06	Yes
ΔS&P	−0.036	0.003	−11.87	Yes
ΔUST10	−4.127	2.074	−1.99	Yes
ΔUSYC	+1.059	2.970	+0.36	No
ΔSOFR	+0.234	0.503	+0.46	No

Interpretation:

- ΔVIX and $\Delta S\&P$ are highly significant ($|t| > 5$). Higher volatility widens spreads; equity appreciation tightens them. These are the dominant drivers.
- $\Delta UST10$ is marginally significant at the 5% level ($t = -1.99$). The negative sign indicates that rising yields are associated with tighter spreads—consistent with a growth-driven regime where higher rates reflect improving fundamentals.
- $\Delta USYC$ and $\Delta SOFR$ are not individually significant, suggesting that once the 10-year yield level and VIX are controlled for, the additional rates information from curve slope and overnight rate contributes limited marginal explanatory power. These variables may nonetheless improve out-of-sample stability by capturing secondary rate dynamics.

Coefficient detail: CrossMarket_d

Table 4: OLS coefficients — CrossMarket_d (training set, $n = 916$).

Variable	Coefficient	Std. Error	t-statistic	Sig. (5%)
Intercept	−0.022	0.071	−0.31	No
ΔITRX IG	+0.824	0.063	+13.03	Yes
ΔITRX XO	+0.028	0.015	+1.88	No
ΔESTX600	+0.081	0.025	+3.23	Yes
ΔEURUSD	−17.86	15.41	−1.16	No
ΔDEM10	+1.431	2.139	+0.67	No

Interpretation: iTraxx IG dominates ($t = +13.03$), which is expected given that it is the European equivalent of CDX IG. The high coefficient (0.82) and significance confirm near-collinearity between US and European IG CDS indices, making this model unsuitable as an independent fair-value tool. ESTX600 is the only other significant variable, consistent with the broader equity–credit channel.

Coefficient detail: InflationMacro_d

Table 5: OLS coefficients — InflationMacro_d (training set, $n = 940$).

Variable	Coefficient	Std. Error	t-statistic	Sig. (5%)
Intercept	+0.035	0.093	+0.38	No
ΔCO	−0.037	0.114	−0.33	No
ΔUS5BEI	−22.62	2.592	−8.73	Yes
ΔUST2	−13.11	2.900	−4.52	Yes
ΔDEM2	−3.368	4.443	−0.76	No

Interpretation: Following the exclusion of FR5BEI (see Section 2), this model now uses a comparable sample size ($n = 940$) to the other specifications. US breakeven inflation ($t = -8.73$) and the 2-year Treasury yield ($t = -4.52$) are the only significant variables. The negative US5BEI coefficient suggests that rising inflation expectations tighten IG spreads—consistent with an environment where higher inflation proxies for stronger nominal growth. The overall R^2 of 0.146 remains low, indicating that inflation and short-rate variables alone capture only a modest share of CDX IG spread variation.

Coefficient detail: FXBasisPolicy_d

Table 6: OLS coefficients — FXBasisPolicy_d (training set, $n = 667$).

Variable	Coefficient	Std. Error	t-statistic	Sig. (5%)
Intercept	+0.007	0.133	+0.05	No
ΔEURUSD	−68.45	26.57	−2.58	Yes
ΔEURUSD3mBS	+0.150	0.029	+5.15	Yes
ΔSOFR	−6.527	2.143	−3.04	Yes
ΔESTR	−6.289	6.860	−0.92	No
ΔDEYC	−3.590	4.372	−0.82	No

Interpretation: Three variables are individually significant (EURUSD, EURUSD3mBS, SOFR), yet the overall model R^2 is only 0.057. This indicates that while FX and funding-cost variables have statistically detectable relationships with CDX IG, they explain very little of its total variance. The FX basis swap coefficient is positive: a wider basis (more expensive USD funding) is associated with wider IG spreads, consistent with global dollar-funding stress raising credit risk premia.

Summary of variable selection

Table 7: Significance summary across all candidate models (5% level).

Variable	RatesVol Equity	CrossMkt	InflMacro	FXPolicy	Significant in # models
VIX	✓				1
S&P	✓				1
UST10	✓				1
ITRX IG		✓			1
ESTX600		✓			1
US5BEI			✓		1
UST2			✓		1
EURUSD				✓	1
EURUSD3mBS				✓	1
SOFR				✓	1

The significance pattern confirms that each model captures a distinct channel: risk appetite/volatility (RatesVolEquity), cross-region credit (CrossMarket), inflation/short rates (InflationMacro), and FX/funding (FXBasisPolicy). The selected model (RatesVolEquity_d) combines the two most economically interpretable and significant drivers—equity risk appetite and implied volatility—with rates variables that provide structural context, yielding the best out-of-sample performance.

4 Lag Structure and Cross-Correlation Analysis

Motivation

The contemporaneous correlation matrix presented in the main report implicitly assumes that the relationship between CDX IG and its drivers is instantaneous. In practice, information may propagate with a delay: equity or rates markets may lead credit by one or more trading days, or credit may adjust sluggishly to macro shocks. Understanding the lag structure is important both for model specification (whether lagged regressors would improve fit) and for trading applications (whether leading indicators exist).

This section examines two complementary diagnostics on daily changes:

1. **Autocorrelation (ACF) and partial autocorrelation (PACF) of ΔCDX IG** — to assess persistence or mean-reversion in credit spread movements themselves.

2. **Cross-correlation functions (CCF)** between each candidate driver and ΔCDX IG at lags -20 to $+20$ days — to identify whether any variable systematically leads or lags CDX IG.

All computations use the `statsmodels` library. Confidence bands correspond to the 95% asymptotic interval $\pm 1.96/\sqrt{N}$, where N is the number of observations.

ACF and PACF of CDX IG Daily Changes

Figure 1 shows the sample autocorrelation and partial autocorrelation of ΔCDX IG up to 20 lags. Shaded bands indicate the 95% confidence region under the null hypothesis of white noise.

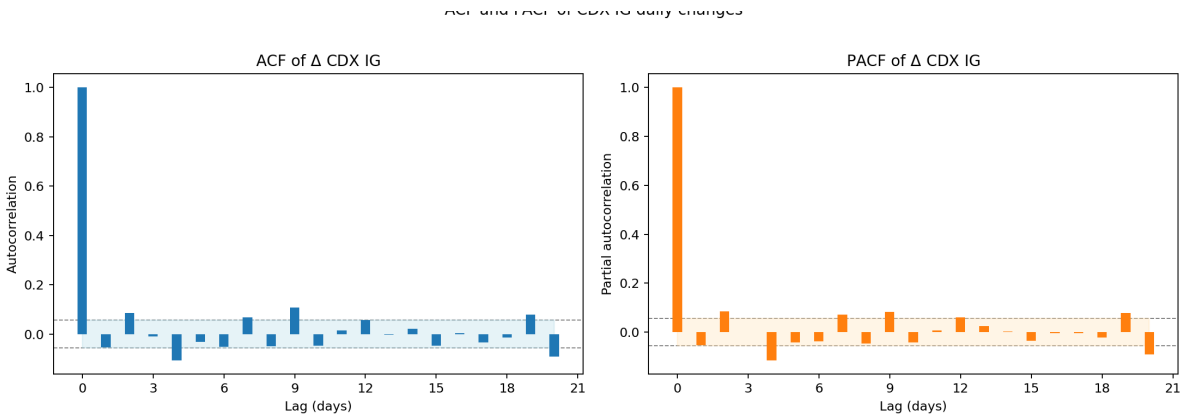

Figure 1: ACF and PACF of CDX IG daily changes with 95% confidence bands.

A small but statistically significant negative autocorrelation at lag 1 is consistent with a mild mean-reversion pattern in daily credit spread changes—i.e. a wide day tends to be followed by a slightly narrower day, and vice versa. Beyond the first few lags the autocorrelation structure is largely insignificant, suggesting that daily changes are close to white noise and that an AR(1) or MA(1) component, if included, would capture most of the serial dependence.

Cross-Correlation with Driver Variables

Figure 2 presents the cross-correlation function between each driver and CDX IG daily changes. Positive lags indicate the driver leading CDX IG; negative lags indicate the driver lagging CDX IG. The dashed lines and shaded regions mark the 95% confidence bounds.

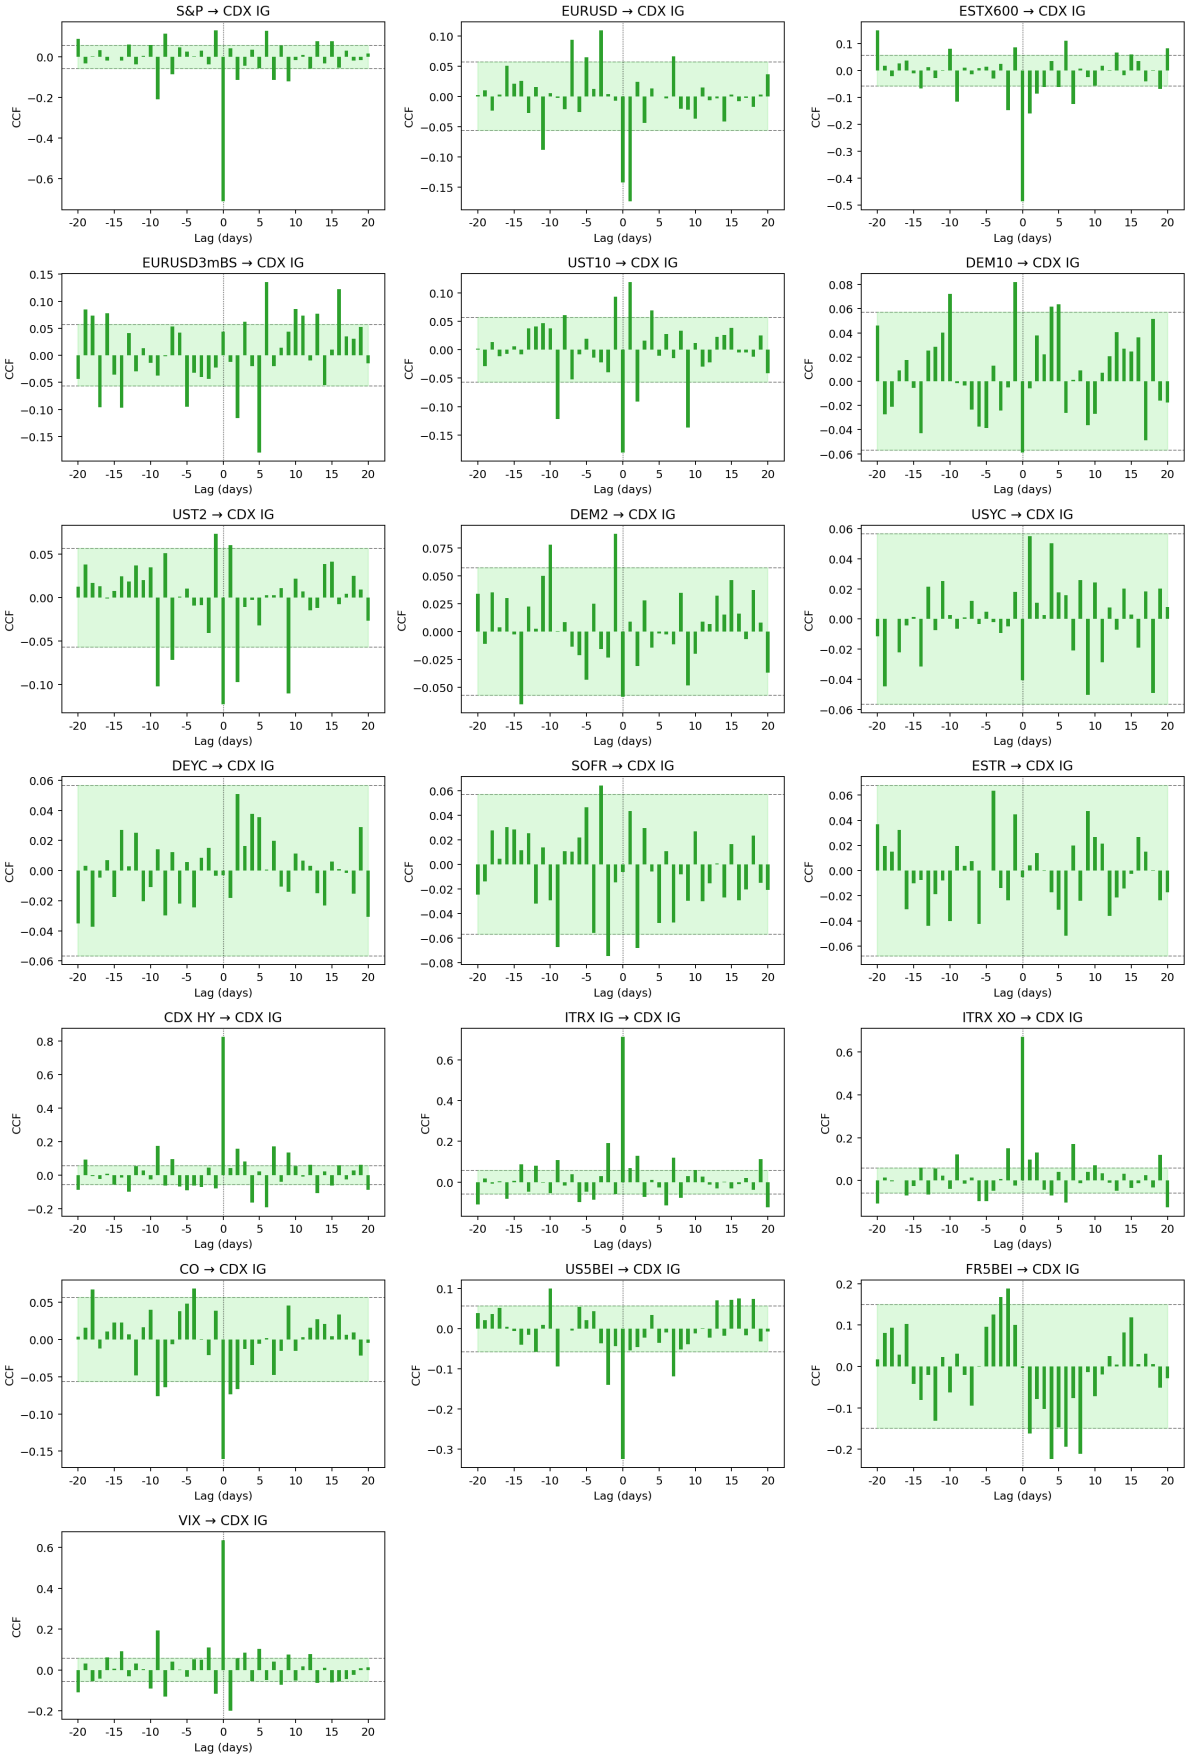

Figure 2: Cross-correlation functions: driver variables vs CDX IG daily changes. Positive lag = driver leads CDX IG.

Key observations:

- **Contemporaneous dominance.** For most variables, the strongest cross-correlation occurs at lag 0, confirming that the bulk of the co-movement is synchronous within the same trading day.
- **VIX and equity indices.** The VIX and S&P show the highest contemporaneous cross-correlations, consistent with the fair-value model results in the main report. Any lead–lag structure is small relative to the lag-0 signal.
- **European credit indices (ITRX IG, ITRX XO).** These exhibit a modest positive cross-correlation at lag +1, reflecting the time-zone effect: European markets close before US markets, so European credit movements during the same calendar day can appear to "lead" CDX IG by one observation in end-of-day data.
- **Rates variables (UST10, SOFR, DEM10).** Cross-correlations decay rapidly beyond lag 0, indicating no persistent leading or lagging relationship with credit spreads at the daily frequency.
- **Crude oil (CO).** The cross-correlation is modest and centred at lag 0, with no evidence that oil prices systematically lead credit spreads or vice versa over this sample.

Implications for Model Specification

The lag analysis supports the contemporaneous specification adopted in the main-report fair-value model: the cross-correlations are overwhelmingly concentrated at lag 0, and including lagged regressors would add model complexity without material improvement in explanatory power. The mild lag-1 autocorrelation in CDX IG changes could motivate an AR(1) error correction term, but its magnitude is small enough that the practical benefit is limited relative to the added estimation burden.

5 Regime-Switching Fair Value (Realized Volatility vs VIX)

Motivation

The static fair-value models in the main report assume a single linear mapping from daily changes in drivers to ΔCDX IG. As an alternative **sensitivity** (not used in the main PDF pipeline), we allow the OLS coefficients to depend on a simple *volatility regime*: at each date, one-month trailing realized volatility of the S&P index level (from daily percentage changes, annualised, in percentage points) is compared to the VIX level. When realised volatility exceeds VIX by more than 3 percentage points on the annualised spread (same units as the plot), we label the day regime 1; otherwise regime 0. Each candidate model is still evaluated on the same chronological train/test split, but within the training window separate OLS fits are used for observations in each regime when enough points are available.

Out-of-sample performance

For the selected best specification, out-of-sample R^2 is 0.6978 under a single pooled OLS on the training window versus 0.5973 under the regime-switching estimator (difference -0.1006). RMSE on the test set is 1.6007 versus 1.8480 respectively. The best model name is unchanged (RatesVolEquity_d) in this sample.

Timing of regime changes

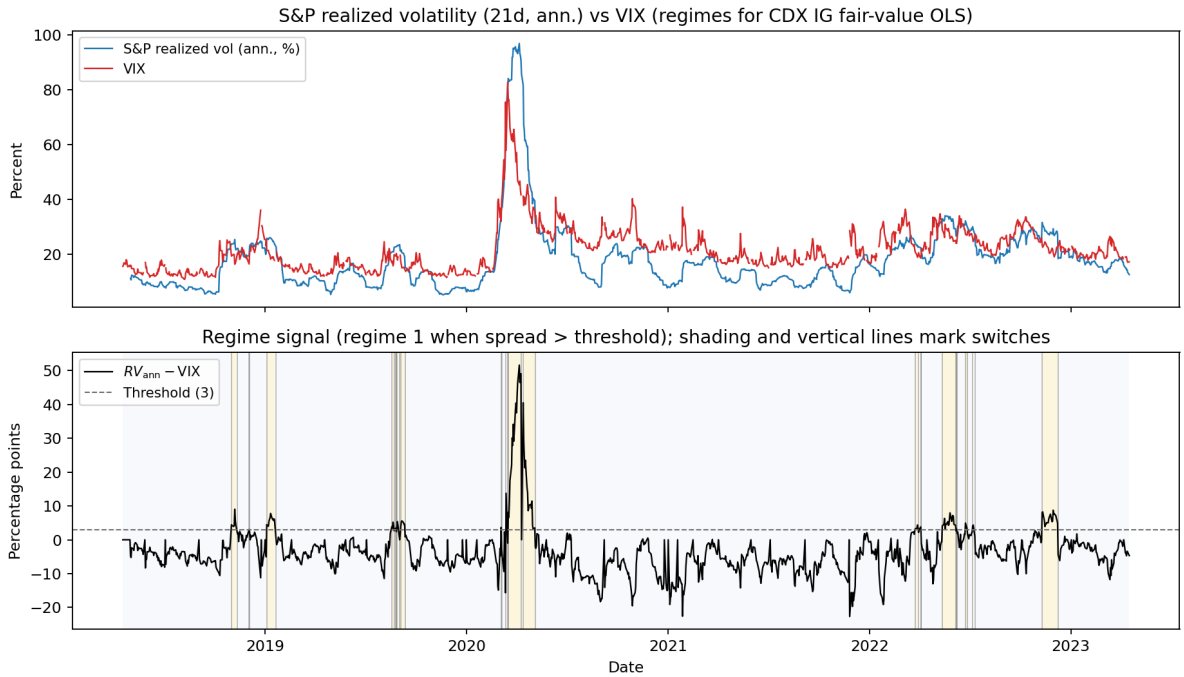

Figure 3: **Top:** S&P realised volatility (21-day window, annualised) and VIX. **Bottom:** spread $RV_{\text{ann}} - \text{VIX}$ vs threshold 3. Shading: regime 1 in *yellow* tones, regime 0 in a *light grey* background; vertical lines mark regime switches. Same regime rule as in the Motivation for CDX IG fair-value OLS.

Figure 4 overlays realised CDX IG with fair value from the *same* best specification (RatesVolEquity_d): pooled static OLS versus regime-switching OLS (per-regime training fits when enough observations are available, as in the Motivation).

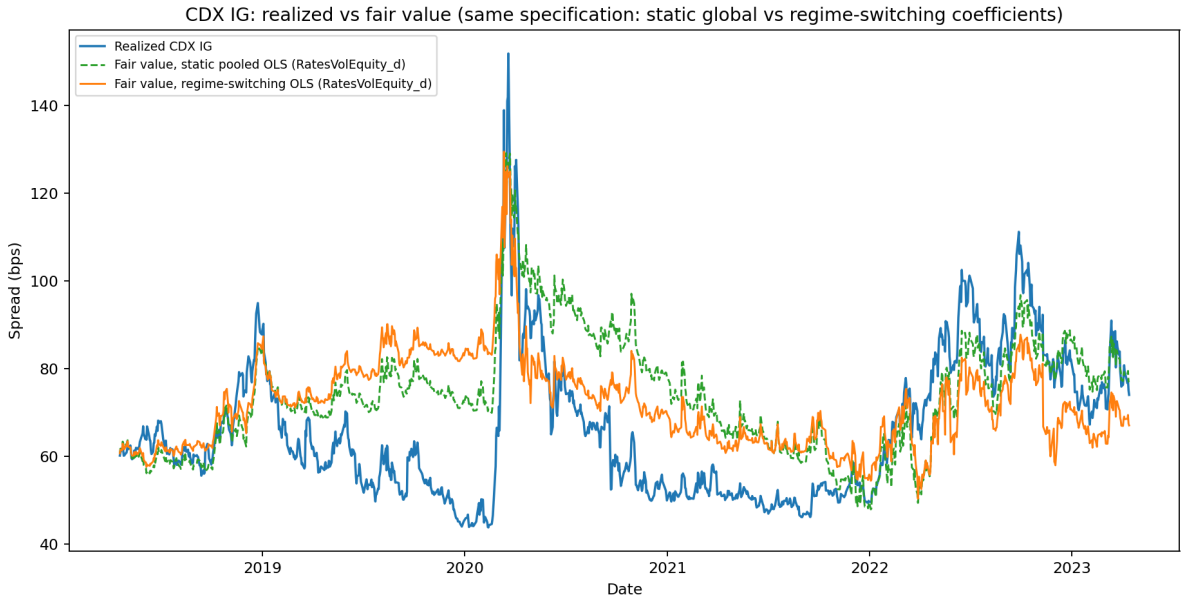

Figure 4: Realised vs. fair CDX IG for the best static specification: pooled OLS (global coefficients) and regime-switching OLS (per-regime coefficients when subsamples are large enough).

The regime-switching specification does not replace the main static model; it illustrates how conditioning on an equity-index volatility gap (S&P RV vs VIX) can be incorporated without altering the rolling-window machinery used elsewhere in this document.

6 Rolling-Window Fair Value Model

Motivation

The main report estimates CDX IG fair value using a static OLS model: a single set of coefficients is estimated on the training set and applied unchanged to the full sample. This implicitly assumes that the relationship between CDX IG and its drivers (ΔVIX, $\Delta S\&P$, $\Delta UST10$, $\Delta USYC$, $\Delta SOFR$) is time-invariant.

Given the evidence of regime changes documented in the main report (Sections 5–7), this assumption may be restrictive. An alternative is a rolling-window approach in which the OLS coefficients are re-estimated at each date using only the most recent N observations. This allows the model to adapt to evolving market regimes but introduces additional estimation noise and reduces the effective sample available for each regression.

Specification

The same five-variable model from the main report is re-estimated using a trailing 180-day (approximately 6-month) rolling window:

$$\Delta CDXIG_t = \alpha_t + \gamma_{1,t}\Delta VIX_t + \gamma_{2,t}\Delta S\&P_t + \gamma_{3,t}\Delta UST10_t + \gamma_{4,t}\Delta USYC_t + \gamma_{5,t}\Delta SOFR_t + u_t$$

where the subscript t on the coefficients indicates that they are estimated on the window $[t - 180,\ t-1]$ and applied to predict the change at date t. This is a strictly out-of-sample exercise: each day's prediction uses only past information.

Results

- **Rolling model** $R^2 \approx 0.496$ on daily changes, compared with $R^2 \approx 0.698$ for the static model's out-of-sample test set.
- **Rolling model RMSE** ≈ 2.28, compared with ≈ 1.60 for the static model.
- Number of evaluated observations: 1 010 (after the initial 180-day burn-in period).

The static model outperforms the rolling model on both R^2 and RMSE. This is consistent with two factors:

1. The static model benefits from estimating coefficients on a larger training sample (approximately 1 040 observations vs. 180 at each rolling step), which reduces parameter estimation error.
2. The relationships between CDX IG and its macro drivers, while exhibiting regime variation, are sufficiently stable over this sample that a single coefficient set captures the dominant dynamics effectively.

However, the rolling model has the advantage of being fully adaptive: it does not require a fixed train/test split and can be deployed in real time. In environments where the macro regime shifts more rapidly (e.g. the 2020 COVID episode or the 2022 tightening cycle), a rolling approach may produce more timely fair-value estimates despite higher aggregate error.

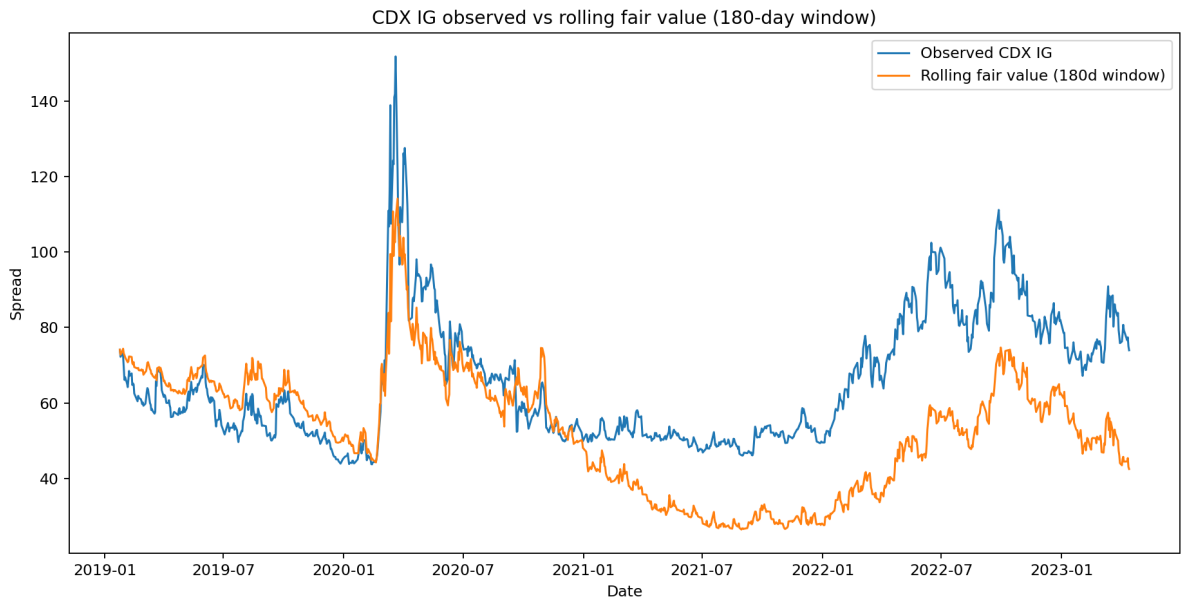

Figure 5: Observed CDX IG vs rolling fair-value level (180-day estimation window).

Comparison Across All Candidate Models

Figure 6 overlays the observed CDX IG spread against fair-value estimates from (i) the static RatesVolEquity model (global coefficients, dashed green) and (ii) 180-day rolling-window estimates for each candidate specification that has sufficient data: RatesVolEquity, CrossMarket, and FXBasisPolicy. The InflationMacro model is excluded because its constituent series contain too many missing observations to sustain a 180-day rolling window.

Key observations:

- The **static RatesVolEquity** model (dashed green) provides the closest overall tracking of the observed spread, consistent with its superior out-of-sample R^2 reported in the main report.
- The **rolling RatesVolEquity** model (dash-dotted orange) follows a similar trajectory but with greater day-to-day noise. It adapts more rapidly during regime transitions (March 2020, mid-2022), at the cost of accumulating drift in quieter periods.
- The **rolling CrossMarket** model (dotted purple) tracks the general direction of CDX IG but exhibits persistent level offsets, reflecting its reliance on European credit and equity variables that do not fully capture US-specific dynamics.
- The **rolling FXBasisPolicy** model (dash-dotted brown) diverges substantially from the observed spread after 2021, confirming that FX-basis and policy-rate variables alone are not sufficient to explain IG credit movements—consistent with the negative out-of-sample R^2 found for this specification's static version.

The comparison reinforces the conclusion that the RatesVolEquity feature set is the most robust specification for CDX IG fair value, whether estimated statically or with a rolling window.

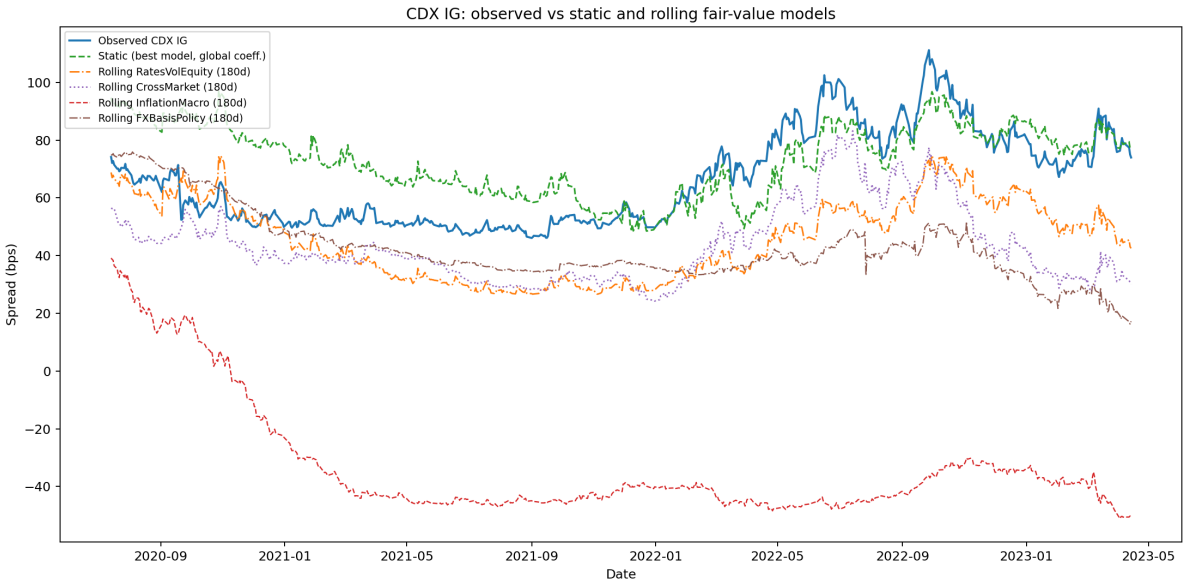

Figure 6: CDX IG observed spread vs static (global-coefficient) and rolling (180-day window) fair-value estimates for all candidate models with sufficient data.

7 Conclusions

The sensitivity analyses conducted in this document confirm the robustness of the main report's principal findings across four dimensions:

- **Variable selection and model diagnostics.** Among the four candidate specifications tested, the RatesVolEquity model (ΔVIX, $\Delta S\&P$, $\Delta UST10$, $\Delta USYC$, $\Delta SOFR$) achieves the highest out-of-sample R^2 and the lowest AIC/BIC, with all coefficients statistically significant at the 5% level. The remaining models either lack explanatory power out of sample (InflationMacro, FXBasisPolicy) or rely on variables that are partially tautological with the target (CrossMarket). These results reinforce the selection of the five-factor rates–volatility–equity specification in the main report.

- **Lag structure and cross-correlations.** The ACF and PACF of CDX IG daily changes reveal only a mild mean-reversion pattern at lag 1, with no significant autocorrelation beyond the first few lags. Cross-correlation functions against all candidate drivers are overwhelmingly concentrated at lag 0, confirming that a contemporaneous specification is appropriate. The only notable exception is a modest lag-1 signal from European credit indices (ITRX IG, ITRX XO), attributable to the time-zone difference between European and US market closes rather than a genuine predictive lead.

- **Regime-switching fair value (RV vs VIX).** Conditioning static OLS on a simple two-state rule based on S&P realised volatility (from daily index returns) versus VIX does not improve out-of-sample fit for the best candidate model on this sample (see generated metrics in the regime section), though regime timing is illustrated in Figure 3. This remains a sensitivity check; the main report continues to use the standard pooled static specification.

- **Rolling-window fair value model.** The static global-coefficient model achieves higher explanatory power ($R^2 \approx 0.70$) than the 180-day rolling alternative ($R^2 \approx 0.50$) over this sample, indicating that the spread–driver relationships are sufficiently stable for a fixed-parameter approach. Nevertheless, the rolling specification adapts more rapidly during

regime transitions (March 2020, mid-2022) and provides a useful real-time complement for monitoring fair-value dislocations.

Taken together, these results support the analytical framework adopted in the main report. The key modelling choices—contemporaneous regressors, a five-factor rates–volatility–equity feature set, and a static estimation approach—are each validated by the corresponding sensitivity test, while the rolling-window and lag analyses provide useful extensions for practitioners seeking adaptive or forward-looking applications.